

mlpnn

November 11, 2021

1 MLPNN Exercise

The target of this assignment is to design and implement bp-MLPNN with at least two different error functions for [MNIST](#) handwritten digit classification.

1.1 Table of Contents

- 1-Packages
- 2-Load the Dataset
- 3-Skeleton of SVM
- 4-Training with Cross-Entropy Loss
- 5-Training with KL-Divergence Loss

1.2 1 - Packages

First import all the packages needed during this assignment.

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn import metrics
from sklearn.datasets import fetch_openml
from tqdm import tqdm

%matplotlib inline

%load_ext autoreload
%autoreload 2

np.random.seed(1) # set a seed so that the results are consistent
```

1.3 2 - Load the Dataset

```
[2]: # Load MNIST dataset via sklearn
data, target = fetch_openml('mnist_784', return_X_y=True)

data = data.to_numpy() / 255.0
target = target.to_numpy().astype(np.uint8)

n_samples = 60000 # MNIST dataset uses 60000 examples for training
```

```

train_img = data[:n_samples, :]
train_label = target[:n_samples]
test_img = data[n_samples:]
test_label = target[n_samples:]

```

```

[3]: # As a sanity check, print out the size of the training and test data
print('Training data shape: ', train_img.shape)
print('Training labels shape: ', train_label.shape)
print('Test data shape: ', test_img.shape)
print('Test labels shape: ', test_label.shape)

```

```

Training data shape: (60000, 784)
Training labels shape: (60000,)
Test data shape: (10000, 784)
Test labels shape: (10000,)

```

1.4 3 - Backbone of MLPNN

```

[4]: # Perceptron
class Perceptron:
    def __init__(self, dim, activation):
        super().__init__()
        self.act = activation
        self.w = np.random.randn(dim, 1) # column vector
        self.b = np.random.rand()

    def __call__(self, x):
        '''forward pass'''
        Z = np.dot(x, self.w) + self.b
        if self.act:
            A = self.act(Z)
            return Z, A
        else:
            return Z

```

```

[5]: def softmax(X):
    X_exp = np.exp(X)
    partition = np.sum(X_exp, axis=1, keepdims=True)
    return X_exp / partition

```

```

[6]: class MLPNN:
    def __init__(self, in_dim, out_dim, n_neuron):
        super().__init__()

        self.layer1 = [Perceptron(in_dim, activation=np.tanh) for _ in
↪range(n_neuron)] # use tanh for activation

```

```

        self.layer2 = [Perceptron(n_neuron, activation=softmax) for _ in
→range(out_dim)] # softmax activation

        # vectorize to accerlate
        self._vectorize()

        # for back-prop
        self.cache = dict()

def __call__(self, x, requires_cache=True):
    '''forward pass'''

    # Input Layer -> Hidden Layer
    layer1_input = x
    Z1 = np.dot(layer1_input, self.W1) + self.b1
    A1 = self._activate(Z1, layer_idx=1)

    if requires_cache:
        self.cache['Z1'] = Z1
        self.cache['A1'] = A1

    # Hidden Layer -> Output Layer
    layer2_input = A1
    Z2 = np.dot(layer2_input, self.W2) + self.b2
    A2 = self._activate(Z2, layer_idx=2) # size*C, probability of each class

    if requires_cache:
        self.cache['Z2'] = Z2
        self.cache['A2'] = A2

    return A2.clip(min=1e-8, max=None)

def _vectorize(self):
    self.W1 = np.hstack([neuron.w for neuron in self.layer1]) # dim *
→n_hidden
    self.b1 = np.array([neuron.b for neuron in self.layer1]) # n_hidden
    self.W2 = np.hstack([neuron.w for neuron in self.layer2]) # n_hidden *
→10
    self.b2 = np.array([neuron.b for neuron in self.layer2]) # 10

    self.act1 = [neuron.act for neuron in self.layer1]
    if all(x==self.act1[0] for x in self.act1): # check if all perceptrons
→are using the same activation
        self.act1 = self.act1[0]

    self.act2 = [neuron.act for neuron in self.layer2]

```

```

        if all(x==self.act2[0] for x in self.act2): # check if all perceptrons
→are using the same activation
            self.act2 = self.act2[0]

        del self.layer1
        del self.layer2

    def _activate(self, feat, layer_idx):
        assert layer_idx in (1, 2) # since we're using two layer neural network

        act = self.act1 if layer_idx == 1 else self.act2

        if callable(act):
            return act(feat)
        elif type(act) is list:
            assert len(act) == feat.shape[1]
            for i in range(feat.shape[1]):
                feat[:, i] = act[i](feat[:, i])
            return feat
        else:
            raise Exception('Internal error')

```

1.5 4 - Training with Cross-Entropy Loss

For Multi-class classification problem, cross-entropy loss is defined as

$$L(\hat{y}, y) = - \sum_k^K y^{(k)} \log \hat{y}^{(k)}$$

where $y^{(k)}$ is 0 or 1, indicating whether class label k is the correct classification.

```

[7]: def to_one_hot(values):
    values = np.squeeze(values)
    n_values = np.max(values) + 1
    return np.eye(n_values)[values].astype(np.uint8)

def compute_loss(output, Y):
    """ Computes cross entropy between two distributions.
    Input: x: iterable of N non-negative values
           y: iterable of N non-negative values
    Returns: scalar
    """
    return -np.sum(np.multiply(to_one_hot(Y), np.log(output)))

[8]: def bp_cross_entropy(network, X, Y, lr):
    n = X.shape[0] # size of dataset

```

```

A1 = network.cache['A1']
A2 = network.cache['A2']

# compute gradients
dZ2 = A2 - to_one_hot(Y)
dW2 = (1 / n) * np.dot(A1.T, dZ2)
db2 = (1 / n) * np.sum(dZ2, axis=0)

dZ1 = np.multiply(np.dot(dZ2, network.W2.T), 1 - np.power(A1, 2))
dW1 = (1 / n) * np.dot(X.T, dZ1)
db1 = (1 / n) * np.sum(dZ1, axis=0)

# update parameters
network.W2 = network.W2 - lr * dW2
network.b2 = network.b2 - lr * db2
network.W1 = network.W1 - lr * dW1
network.b1 = network.b1 - lr * db1

```

```

[9]: network = MLPNN(in_dim=784, out_dim=10, n_neuron=1000)
learning_rate = 1e-1
epochs = 2000

training_loss = list()
testing_loss = list()

for epoch in tqdm(range(1, epochs+1)):
    out = network(train_img)
    loss = compute_loss(out, train_label)
    training_loss.append(loss)
    bp_cross_entropy(network, train_img, train_label, learning_rate)

    if epoch % 200 == 0:
        print('[epoch %d] loss: %.4f' % (epoch, loss))

    # Evaluate on testing set
    out = network(test_img, requires_cache=False)
    loss = compute_loss(out, test_label)
    testing_loss.append(loss)

# Plot loss curve
p1, = plt.plot(training_loss)
p2, = plt.plot(testing_loss)
plt.legend(handles=[p1, p2], labels=['train', 'test'], loc='best')
plt.xlabel('Epochs')
plt.ylabel('Cross-Entropy Loss')

```

10%|


| 199/2000 [18:11<2:48:54, 5.63s/it]

```

[epoch 200] loss: 140015.9757
20%|                               | 399/2000 [36:09<2:20:03,  5.25s/it]
[epoch 400] loss: 98761.8062
30%|                               | 599/2000 [53:27<1:59:15,  5.11s/it]
[epoch 600] loss: 79200.0926
40%|                               | 799/2000 [1:10:44<1:51:25,  5.57s/it]
[epoch 800] loss: 66535.5184
50%|                               | 999/2000 [1:28:43<1:29:44,  5.38s/it]
[epoch 1000] loss: 57299.7318
60%|                               | 1199/2000 [1:46:49<1:07:26,  5.05s/it]
[epoch 1200] loss: 50137.3821
70%|                               | 1399/2000 [2:03:41<50:41,  5.06s/it]
[epoch 1400] loss: 44356.6473
80%|                               | 1599/2000 [2:20:32<33:44,  5.05s/it]
[epoch 1600] loss: 39598.9911
90%|                               | 1799/2000 [2:37:24<16:55,  5.05s/it]
[epoch 1800] loss: 35571.1231
100%|                             | 1999/2000 [2:54:16<00:05,  5.06s/it]
[epoch 2000] loss: 32051.5763
100%|                             | 2000/2000 [2:54:21<00:00,  5.23s/it]

```

```
[9]: Text(0, 0.5, 'Cross-Entropy Loss')
```



```
[10]: def evaluate(network, X, Y):
    pred = np.argmax(network(X), axis=1).reshape(-1)
    cm = metrics.confusion_matrix(Y, pred)
    print("Confusion matrix:\n%s" % cm)
    print("Accuracy={}".format(metrics.accuracy_score(Y, pred)))

    evaluate(network, test_img, test_label)
```

Confusion matrix:

```
[[ 914   0   6   3   2  23  13   6  11   2]
 [   0 1098   8   6   1   3   3   2  13   1]
 [  11   9  898  29  11   5  14  11  37   7]
 [   4   4  12  890   1  44   1  11  33  10]
 [   3   3  10   1 866   3  14  18  12  52]
 [  12   3   9  41  11 744  13   8  43   8]
 [  17   3  18   4  17  15 876   2   4   2]
 [   4   7  27  12  11   2   1 919   7  38]
 [  13   6  21  35  11  41  14  14 801  18]
 [  10   5   1  14  58   6   4  33  18 860]]
```

Accuracy=0.8866

1.6 5 - Training with KL-Divergence Loss

Kullback Leibler Divergence, or KL Divergence for short, is a measure of how one probability distribution differs from a baseline distribution. For Multi-class classification problem, KL-Divergence

loss is defined as

$$L(\hat{y}, y) = \sum_k^K y^{(k)} (\log \hat{y}^{(k)} - \log y^{(k)})$$

where $y^{(k)}$ is 0 or 1, indicating whether class label k is the correct classification.

```
[11]: def compute_loss(output, Y):  
    y = to_one_hot(Y)  
    return -np.sum(np.where(y != 0, np.multiply(y, np.log(output / y)), 0))
```

```
[12]: # Note that code is same with back prop in cross-entropy  
# since both losses has same derivative w.r.t y hat  
def bp_kl(network, X, Y, lr):  
    n = X.shape[0] # size of dataset  
  
    A1 = network.cache['A1']  
    A2 = network.cache['A2']  
  
    # compute gradients  
    dZ2 = A2 - to_one_hot(Y)  
    dW2 = (1 / n) * np.dot(A1.T, dZ2)  
    db2 = (1 / n) * np.sum(dZ2, axis=0)  
  
    dZ1 = np.multiply(np.dot(dZ2, network.W2.T), 1 - np.power(A1, 2))  
    dW1 = (1 / n) * np.dot(X.T, dZ1)  
    db1 = (1 / n) * np.sum(dZ1, axis=0)  
  
    # update parameters  
    network.W2 = network.W2 - lr * dW2  
    network.b2 = network.b2 - lr * db2  
    network.W1 = network.W1 - lr * dW1  
    network.b1 = network.b1 - lr * db1
```

```
[13]: network = MLPNN(in_dim=784, out_dim=10, n_neuron=1000)  
learning_rate = 1e-1  
epochs = 2000  
  
training_loss = list()  
testing_loss = list()  
  
for epoch in tqdm(range(1, epochs+1)):  
    out = network(train_img)  
    loss = compute_loss(out, train_label)  
    training_loss.append(loss)  
    bp_kl(network, train_img, train_label, learning_rate)
```



```

if epoch % 200 == 0:
    print('[epoch %d] loss: %.4f' % (epoch, loss))

# Evaluate on testing set
out = network(test_img, requires_cache=False)
loss = compute_loss(out, test_label)
testing_loss.append(loss)

# Plot loss curve
p1, = plt.plot(training_loss)
p2, = plt.plot(testing_loss)
plt.legend(handles=[p1, p2], labels=['train', 'test'], loc='best')
plt.xlabel('Epochs')
plt.ylabel('KL-Divergence Loss')

```

0%| | 0/2000 [00:00<?,
?it/s]/tmp/ipykernel_184119/221692360.py:3: RuntimeWarning: divide by zero
encountered in true_divide
return np.sum(np.where(y != 0, np.multiply(y, np.log(output / y)), 0))
/tmp/ipykernel_184119/221692360.py:3: RuntimeWarning: invalid value encountered
in multiply
return np.sum(np.where(y != 0, np.multiply(y, np.log(output / y)), 0))
10%| | 199/2000 [16:46<2:31:45, 5.06s/it]
[epoch 200] loss: -143104.7662
20%| | 399/2000 [33:38<2:15:24, 5.07s/it]
[epoch 400] loss: -101351.5404
30%| | 599/2000 [50:30<1:58:13, 5.06s/it]
[epoch 600] loss: -81278.7056
40%| | 799/2000 [1:07:21<1:41:17, 5.06s/it]
[epoch 800] loss: -67875.5278
50%| | 999/2000 [1:24:13<1:24:22, 5.06s/it]
[epoch 1000] loss: -58113.3984
60%| | 1199/2000 [1:41:05<1:07:28, 5.05s/it]
[epoch 1200] loss: -50524.0959
70%| | 1399/2000 [1:58:33<54:21, 5.43s/it]
[epoch 1400] loss: -44438.4743
80%| | 1599/2000 [2:16:50<36:43, 5.50s/it]
[epoch 1600] loss: -39468.1728
90%| | 1799/2000 [2:34:03<17:44, 5.29s/it]


```
[epoch 1800] loss: -35320.9407
```

```
100%|          | 1999/2000 [2:51:22<00:05, 5.06s/it]
```

```
[epoch 2000] loss: -31768.7917
```

```
100%|          | 2000/2000 [2:51:27<00:00, 5.14s/it]
```

```
[13]: Text(0, 0.5, 'KL-Divergence Loss')
```



```
[14]: evaluate(network, test_img, test_label)
```

Confusion matrix:

```
[[ 920   1  11   6   1  17  15   2   6   1]
 [   0 1091   6   6   1   4   5   4  17   1]
 [  17   7 906  22   9  12  13  10  31   5]
 [   4   4  16 893   2  33   2  16  27  13]
 [   7   1  12   5 864   3  19   5  14  52]
 [  10   5  13  34  13 751  16   7  31  12]
 [  13   3  14   3  16  25 868   0  15   1]
 [   4   9  31  14  13   8   1 897   3  48]
 [   9   6  16  33  13  37  12  10 819  19]
 [   5   8   6   9  51  11   3  33  24 859]]
```

Accuracy=0.8868